

Berlin G

berlinmivi@gmail.com | linkedin.com/in/berlin-g | github.com/berlinmivi

PROFILE

Machine Learning Engineer with 3+ years of hands-on experience building and productionizing ML models across NLP, forecasting, and optimization use cases. Expertise in Python, classical ML algorithms, and deploying models via APIs, Docker, and AWS in product companies.

KEY SKILLS

Machine Learning

Linear Regression, Logistic Regression, Decision Trees, Gradient Boosting (LightGBM), Time-Series Forecasting, Recommendation Systems, Price Optimization, Feature Engineering

NLP & AI

Text Generation, Text Classification, Embeddings, Similarity Models, User Profiling, Feature Extraction, Transformers, Langchain, RAG Systems

Python & Libraries

Scikit-learn, TensorFlow, Pandas, NumPy, Matplotlib, Prophet, Pulp, LightGBM, SHAP

Productionization

FastAPI, Docker, CI/CD, REST APIs, AWS (Bedrock, Lambda, EC2, S3), Nginx, MLflow, Airflow

Data & Infrastructure

MySQL, PostgreSQL, Redis, Qdrant, ChromaDB, Git, GitHub

EDUCATION

IIT Bombay

M.Tech. Aerospace Engineering

CGPA: 8.01/10 | 2020-2022

Key Course: Introduction to Machine Learning

TA: Mentored students in ML algorithms

MIT Chennai

B.E. Aeronautical Engineering

CGPA: 8.8/10 | 2015-2019

Key Course: Computing Techniques

CERTIFICATIONS

- Introduction to Statistics - Stanford
- Credit Risk Modeling in Python
- Supply Chain Analytics - Rutgers
- ML with Python

ACHIEVEMENTS

- AIR 81 GATE Aerospace 2020

WORK EXPERIENCE

Althi Solutions | ML Engineer [Jul 2024 - Present]

RAG-Based Conversational System

- Engineered real-time NLP assistant using FastAPI, AWS Bedrock & Qdrant with text generation and embeddings; Implemented semantic chunking, similarity models & feature extraction from unstructured text.
- Productionized with Docker on AWS EC2, SNS-triggered webhook for updates, and S3 knowledge base; Integrated Grafana logging with REST API interface.

AML Compliance Assistant

- Built fraud detection and risk assessment agent using Langgraph & AWS Bedrock; Integrated MySQL for transaction pattern analysis, anomaly detection & risk profiling.
- Productionized with FastAPI, Docker, Nginx load balancing, and CI/CD pipeline on AWS; Implemented ModSecurity firewall and Redis authentication.

CloudWatch Triage BOT

- Developed text classification system using Langchain React Agent; Applied NLP for pattern recognition, feature extraction from unstructured logs & automated incident categorization.
- Productionized with Docker and Ollama; Built Streamlit interface for conversational log access.

Ikigai | ML Engineer [Jul 2022 - Jun 2024]

Inventory Optimization - Taager

- Implemented time-series forecasting using Prophet across 20K+ SKUs; Applied feature engineering for seasonal variations, promotions & geographic behavior.
- Designed price optimization and procurement recommendation system using Pulp linear programming with gradient-based optimization; Achieved 10% GMV increase & 9M AED savings.

Episodic Forecasting - Spirit Halloween

- Built forecasting system for 10K+ SKUs using gradient boosting (LightGBM), decision trees, and Prophet; Applied PCA and feature engineering with time-series cross-validation.
- Achieved WMAPE below 23% and \$4M savings; Utilized SHAP values for model interpretability and feature identification.

Time-Series Forecasting Q&A Model

- Developed RAG system using OpenAI API with embeddings, similarity models & text generation; Built knowledge base using ChromaDB vector database for feature extraction from unstructured text.
- Productionized with Docker and Streamlit UI; Implemented REST API interface for portable deployments.

PROJECTS

Response Assistant

- Fine-tuned DeepSeek-R1 using LORA for text generation with user profiling; Implemented feature extraction from unstructured text to mimic communication style and tone.

Airflow ETL Pipeline

- Designed end-to-end ETL pipeline with Airflow orchestration; Extracted via REST API, transformed with Pandas, loaded into PostgreSQL; Demonstrated CI/CD best practices.